
University of Portsmouth
BSc (Hons) Computer Science
Second Year

Data Structures and Algorithms (DSALG)
M21270
September 2023 - January 2024
20 Credits

Thomas Boxall
up2108121@myport.ac.uk

Contents

1	Async lecture - Introduction to Data Structures and ADT (2023-09-30)	2
2	Async lecture - Tools of the Trade I: Efficiency & BigO (2023-09-30)	5
3	Async lecture - Iterative Algorithms and Efficiency (2023-10-05)	7
4	Async lecture - Tools of the Trade II: Recursion (2023-10-05)	10
5	Async lecture - Recursive Algorithms & Efficiency (2023-10-11)	12
6	Async lecture - Linked Lists (2023-10-19)	14
7	Async lecture - Hierarchical Data Structures (2023-11-02)	16
8	Async lecture - Self Balancing Trees (2023-11-09)	20
9	Async Lecture - Self-Organising Trees (2023-11-09)	25

Page 1

Async lecture - Introduction to Data Structures and ADT

📅 2023-09-30



1.1 Data Structures

A *Data Structure* is a way to store and organise data in order to facilitate access and modification. There is no single data structure which is perfect for every application - we need to choose the best for whatever we are creating.

There are two parts to a data structure: a collection of elements, each of which is either a data type of another data structure; and a set of associations or relationships (the structure) involving the collection of elements.

1.1.1 Classification of Data Structures

Data structures can be classified based on their predecessor and successor.

Name	Predecessor	Successor	Examples
Linear	unique	unique	stack, queue
Hierarchical	unique	many	family tree, management structure
Graph	many	many	railway map, social network
Set Structure	no	no	DSALG class

Table 1.1: Classifications of Data Structures

1.1.2 Choosing the right Data Structure

When choosing a data structure, it is important to analyse the problem, determine the basic operations needed and select the most efficient data structure. Choosing the right data structure will make the operations simple & efficient and choosing the wrong data structure will make your operations cumbersome and inefficient.

1.1.3 CRUD

CRUD Operations: Create, Read, Update and Delete are the basic operations which all data structures must be able to do. It is common for a data structure to use a different name to refer to the operation, however.

1.2 Abstract Data Type

An Abstract Data Type (ADT) is a collection of data and associated methods stored as a single module. The data within an ADT cannot be accessed directly, it must be accessed indirectly through its methods. An ADT consists of: the data structure itself; methods to access the data structure; methods to modify the data structure; and internal methods (which are not accessible from outside the ADT).

1.3 Algorithms

An *Algorithm* is any well-defined computational procedure that takes some data, or set of data as input and produces some data or set of data as output. It is the sequence of computational steps which are gone through that transforms the input into the output which is the algorithm. The algorithm must process the data efficiently (both in terms of time and space).

1.3.1 Classifications of Algorithms

There are a number of different classifications of algorithms - four are shown below. Definitions are from *National Institute of Standards and Technology - Dictionary of Algorithms and Data Structures*

1.3.1.1 Brute-Force Algorithm

An algorithm that inefficiently solves a problem, often by trying every one of a wide range of possible solutions. E.g. exhaustive search.

1.3.1.2 Divide & Conquer Algorithm

An algorithm which solves a problem either directly because that instance is easy (typically, this would be because the instance is small) or by dividing it into two or more smaller instances. Each of these smaller instances is recursively solved, and the solutions are combined to produce a solution for the original instance.

1.3.1.3 Backtracking Algorithms

An algorithm that finds a solution by trying one of several choices. If the choice proves incorrect, computation backtracks or restarts at the point of choice and tries another choice. It is often convenient to maintain choice points and alternate choices using recursion.

1.3.1.4 Greedy Algorithms

An algorithm that always takes the best immediate, or local, solution while finding an overall answer. Greedy algorithms find the overall, or globally, optimal solution for some optimisation problems, but may find less-than optimal solutions for some instances of other problems.

1.4 Stack Abstract Data Type

A *stack* is a collection of objects where only the most recently inserted object (**top**) can be removed at any time. A stack is linear and operates in Last In First Out (LIFO).

Stacks must support the following operations.

push add an item to the top of the stack

pop remove an item from the top of the stack

peek examine item at the top of the stack

empty determine if the stack is empty

full determine if the stack is full

An application using a stack will expect the ADT to throw an exception if: a push operation is requested on a full stack; or a push / pop operation is requested on an empty stack. The stack manages its own storage, therefore an application which uses a stack is not concerned with how the storage used by the stack is managed. In general, an ADT is not interested in the application using the ADT.

Stacks can be implemented using a static array and have a number of uses, including: matching brackets in arithmetic expressions; recursive algorithms; and evaluating arithmetic expressions.

1.5 Queue Abstract Data Type

A Queue is a collection of objects organised such that the first object to be stored in the queue is the first to be removed and so on. It is a linear data structure with elements - inserted at one end (the tail) and removed from the other (the head). A queue operates in First In First Out (FIFO).

Queues must support the following operations:

enqueue add item to the queue's tail

dequeue remove item from the queue's head

full check if queue is full (therefore total number of elements exceeds max capacity)

empty check if queue is empty

first check the first element (the head) in the queue

1.5.1 Implementations of a Queue

There are three different implementations of a queue, all of which can use a static array.

1.5.1.1 Fixed Head

The head of the queue is fixed, this means it will always be in index 0 of the static array. When an element is dequeued, the rest of the elements in the queue must be shuffled along so that the new head is in index 0. This is extremely time inefficient for large queues however quite space efficient as there won't be "dead space" at one end of the queue.

1.5.1.2 Mobile Head

The head of the queue is mobile, this means it can be in any index of the array. When an element is dequeued, the rest of the elements in the queue stay where they are and the head pointer is updated to represent the new head's index. This is more time efficient however not space efficient as you may end up with a lot of "dead space" where the head of the queue used to be.

1.5.1.3 Circular Queue

The queue is circular in a logical, not physical, way. This means the head and tail can be anywhere in the array. If the head is not at the start of the array and the space is needed, the tail will loop around to use the space at the start of the static array.

Page 2

Async lecture - Tools of the Trade I: Efficiency & BigO

 2023-09-30





BigO Notation is used to define the efficiency of an algorithm quantitatively. This is a very helpful tool to have when designing algorithms as it allows us to compare multiple algorithms to and understand which is the best one to use.

Name	Notation	Description
Constant	$O(1)$	Algorithm always executes in the same amount of time regardless of the size of the dataset.
Logarithmic	$O(\log_n)$	Algorithm which halves the dataset with each pass, efficient with large datasets, increases execution time at a slower rate than that at which the dataset size increases.
Linear	$O(n)$	Algorithm whose performance declines as the data set grows, reduces efficiency with increasingly large dataset.
Loglinear	$O(N \log_n)$	Algorithm that divides a dataset but can be solved using concurrency on independent divided lists.
Polynomial	$O(N^2)$	Algorithm whose performance is proportional to the size of the dataset, efficiency significantly reduces with increasingly large datasets.
Exponential	$O(2^n)$	Algorithm that doubles with each addition to the dataset in each pass, very inefficient.

Table 2.1: BigO Notation complexity values, listed best to worse

BigO doesn't look at the exact number of operations, it looks at when the size of a problem approximates infinity therefore two very similar algorithms which are slightly different may have the same Big O despite one having double the number of operations.

2.1 Calculating The BigO Value

1. Determine the basic operations (including: assignment, multiplication, addition, subtraction, division, etc)
2. Count how many basic operations there are in the algorithm (some basic algebraic addition required here!)
3. Convert the total number of operations to BigO (done by: ignoring the less dominant terms; and ignoring the constant coefficient - i.e. $2n + 1$ becomes n).

Page 3

Async lecture - Iterative Algorithms and Efficiency

📅 2023-10-05



3.1 Searching Algorithms

When searching a dataset to find the item you are looking for, the aim of an efficient searching algorithm is to exclude elements to reduce the searching space after each comparison.

3.1.1 Sorted vs Unsorted Data

Sorted data, when the data is in either ascending or descending order, makes it easier to locate the item you are searching for. Depending on the algorithm, when using sorted data, we can exclude more elements from the list to search.

Unsorted data, when the data is in a random order, makes it harder to locate the item you are searching for. This is because you have to examine (in the worst case) every element in the array you are searching to realise that the item you are searching for is not in the array.

3.1.2 Sequential (Linear) Search

1. Start at the beginning
2. Check every element of the array in turn until item located or the end of the array reached (therefore item not located)

This search excludes one data item at a time, which is not great. It works on sorted and unsorted data.

The best case BigO is $O(1)$. This is the case where the first item we examine is the item we want to find.

The worst case BigO is $O(n)$ when the searched value is the last element in the array or not in the array, n comparisons are required.

The average case BigO is $O(\frac{n}{2})$ because if the data is distributed randomly, each element has an equal chance to be the one searched for.

Overall, the BigO for a sequential search is $O(n)$.

3.1.3 Binary Search

1. Check the middle element of the array

2. If not found, work out which half of the array the item could be located in and exclude the half which it won't be included in
3. Repeat by searching the half array which may contain the required item by examining the middle element and eliminating half the array
4. Process repeated on the halved array until either find the item or determine that it doesn't exist.

A binary search only works on a sorted array.

The best case BigO is $O(1)$. This is achieved when the value we are searching for is the middle element, hence it is found on the first comparison.

The worst case BigO is $O(\log_2 n)$, which is achieved when the value we are searching for is not found in the array. We therefore need $1 + \log_2 n$ comparisons which gets converted into BigO.

The average case is $O(\log_2 n)$, if the data is distributed randomly - each element has an equal chance to be the one searched for.

The binary search algorithm has an additional cost - it requires a sorted sequence of data items, which incurs cost if the data is not already sorted.

3.2 Sorting Algorithms

A sorting algorithm aims to make comparisons between data items and swap them according to the desired order of the items. All sorting algorithms will involve two basic steps:

1. Compare items
2. Swap elements

3.2.1 Selection Sort

A selection sort works by sorting the array one item at a time. It divides the list into two parts. The sorted section (on the left) is initially empty and the unsorted section (on the right) is initially the complete unsorted list. The smallest (or largest, depending on if we want ascending or descending) is selected from the unsorted array and swapped with the left-most element from the unsorted section. The item is now in the correct final position within the ascending / descending order. The sorted part has increased in size by 1. The process is iterated on the unsorted part of the array until all items have been considered, the array will now be sorted.

The best case BigO is $O(n^2)$, when the array is already sorted hence each element only has to be compared to its direct neighbours to establish this.

The worst case is $O(n^2)$.

Selection sorts are unsuitable for large data sets.

3.2.2 Bubble Sort

A bubble sort works by repeatedly passing through the list, swapping adjacent items if they are in the wrong order. If sorting into ascending (or descending) order, the largest (or smallest) item will be bubbled to the end of the array, hence the name. The process is iterated on the whole list until all items have been considered, therefore the whole list will now be sorted.

Without using a flag, the best case BigO is $O(n^2)$. The worst case BigO is also $O(n^2)$.

When a flag (variable which denotes when the array is sorted, enabling sorting to stop as soon as array detected as sorted) is used, the best case BigO is $O(n)$ and the worst case BigO is $O(n^2)$.

With or without the flag, the algorithm is very slow for large datasets.

3.2.3 Insertion Sort

An insertion sort works by sorting the list one item at a time. The list is split into a sorted and non-sorted part. With each iteration, the next element waiting to be sorted (in the unsorted part) is inserted in its correct location within the sorted list. The process iteration on whole list until all items have been considered; when completed, the list will now be sorted.

The best case BigO is $O(n)$, when the data is already sorted. The worst case BigO is $O(n^2)$.

This algorithm is efficient for sorting nearly-sorted lists.

Page 4

Async lecture - Tools of the Trade II: Recursion

📅 2023-10-05



4.1 Introduction

Computers are better at repeating themselves than humans are. Humans get bored, computers don't.

Iteration is explicit repetition of code. We often use for and while loops to repeat sections of code, both of which use control variables to control the repetitions.

Recursion is an alternative method of repeating code, where the code is repeated implicitly. Recursion occurs when a function or method calls itself.

4.2 Recursion

Recursion is a technique whereby a problem is expressed as sub-problems in a similar / same form to the original problem but smaller in scope. The sub-problems only differ in input or size. Shown below is an example of a recursive algorithm, that never ends.

```
recursive_print_example(i){  
    print(i)  
    recursive_print_example(i+1)  
}
```

Recursion is applied to problems where: a solution is easy to specify for certain conditions (the stopping case, which is required or the program will continue indefinitely); and rules for proceeding to a new state which is either a stopping case or eventually leads to a stopping case (recursive steps) are identified.

Shown below is an actual recursive algorithm with its output:

```
cheers(int times){  
    print("hip")  
    if (times > 0){  
        cheers(times - 1)  
    }  
    print("hooray")  
}  
//outputs: hip hip hip hooray hooray hooray
```

4.3 Pitfalls of Recursion

Recursion must always be used with care and understanding. It is possible to write compact and elegant recursive programs that fail spectacularly at runtime. The main pitfalls of recursive algorithms are as follows:

- Forgetting the stopping case
- Failure to reach a stopping case
- Excessive use of space
- Excessive repeated computations

There are some times, when you really shouldn't use recursions:

- When the algorithm / data structure is not naturally suited to recursion
- When the recursive solution is not shorter and understandable than the linear solution
- When the recursive solution doesn't run in acceptable time limits and / or space limits
- When the intermediate states of the algorithm don't pass the same data to / from them.

Recursion to be continued next week.

Page 5

Async lecture - Recursive Algorithms & Efficiency

📅 2023-10-11



The BigO value of a recursive algorithm is generally way worse than that of a linear algorithm which is performing the same function. Both the time and space efficiency is worse - as it will involve more operations and take up considerably more memory due to the number of function calls made.

5.1 Merge Sort

The *Merge Sort* is an example of a divide and conquer algorithm. This means it works by dividing the bigger problem into smaller problems which can then be solved independently, making it easier to program. The smaller problem's solutions will then be combined back together to make the overall solution. The outline of how it works is shown below:

1. Divide the array of data to be sorted into two equal halves until we have one element in each group. Note: if there is only one element in the list, it must be sorted (this is the stopping case of the recursive algorithm).
2. Merge the smaller arrays in to one large sorted array, merging one array with another array at a time. Repeat this until the entire array is sorted.

5.2 QuickSort

This is another example of a divide and conquer algorithm. The outline of how it works is shown below.

1. Pick a pivot value (this will be an item in the array)
2. Put all values less than the pivot into one array and all those bigger than the pivot into another. Recursively repeat this until you end up with just one element in the arrays at the end (these will be the biggest and smallest elements in the array).
3. Merge back together as the pivot values will now be in the correct order.

The best case BigO is $O(n \log_2 n)$ which is achieved when the pivot splits into two equally sized parts each time.

The worst case BigO is $O(n^2)$ where the pivot is split into very unequal sizes.

The average case BigO is $O(n \log_2 n)$.

5.3 Backtracking

Backtracking is a type of algorithm which continually searches for a solution by constructing partial solutions, using the correct parts of the solution to be the starting point for the next solution attempt. It is a trial and error algorithm.

When the algorithm finds an incorrect partial solution, it returns “up” a layer to where it had a closer match to the correct solution and tries another branch. The algorithm is complete when either all possibilities have been exhausted or the item has been found.

Page 6

Async lecture - Linked Lists

 2023-10-19





6.1 Linear Data Structures: a recap

A *linear data structure* is a collection of nodes (elements) where each node has a unique predecessor and unique successor. We saw in Workshop 1 that there are also other types of data structure where there can be many predecessors and many successors.

6.1.1 Example: Static Array

A static array is a static sized structure, which is different from a dynamic sized structure for example the ArrayList in Java. The advantage to using a Static Array is that there is faster access to each node providing as the node is known, with a BigO of $O(1)$ to visit `arr[i]`. The disadvantages of using a static array is that its fixed size means it cannot be easily extended (as there has not been enough space allocated in memory) and it cannot be reduced (as there would then be wasted space in memory); static arrays can also be expensive to maintain in terms of time, especially for a group of sorted elements.

CRUD Operation	Data type	Efficiency
Searching	Unsorted Data (linear search)	$O(n)$
	Sorted Data (binary search)	$O(\log_2 n)$
Insertion	Unsorted data	$O(1)$
	Sorted data	$O(n)$
Deletion	Unsorted data	$O(1)$
	Sorted data	$O(n)$

Table 6.1: Efficiency of CRUD operations for a static array

The increased BigO value for the sorted array is due to the fact that the current elements have to be shifted within the array to maintain the sorted status, which obviously increases the number of operations which need to be performed.

6.2 Linked List

A linked list is a collection of objects called *nodes*. Every node has 2 components - information to be stored (the data) and the reference to the next node in the list (the link). A linked list is s dynamic data structure, which means the number of nodes in the list is not fixed and it can grow or shrink on demand.

6.2.1 Singly Linked List

This is the most common type of linked list whereby each node is only linked to one predecessor and one successor - exceptions being the head which is only linked to one successor and the tail which is only linked to the predecessor.

Advantages of using the Singly Linked List is that it is easily extendable or reduced to fit the data needing to be stored in it. Once the item is located - it has efficiency $O(1)$ to insert or delete the item.

The significant drawback of a SLL is that it does not allow direct access to individual items; to find an individual item - you have to start at the head and follow references until you find the item you're looking for, which has the efficiency $O(n)$. Another disadvantage is that it uses more memory compared to a static array as the reference to the subsequent node also has to be stored.

6.3 Other Types of Linked Lists

Alongside the SLL, there are other types of Linked Lists which we may come across.

6.3.1 Singly Linked List with Dummy Nodes

A significant issue with the SLL is that the head item has to point to something. The idea of using a dummy node is that this is the first node in the list which then means the head element can always point to something, even if it is an empty node which makes the code more efficient.

6.3.2 Circular Linked List

In this linked list, the tail node points to the first node after the head (which is the first node to contain data). This is useful for a list which undergoes a large number of traversals.

6.3.3 Doubly Linked List

In this linked list, each node is linked to both its predecessor and successor in a way such that the list can be traversed in either direction.

6.3.4 SkipList

A SkipList is an extension of an ordered SLL with a number of additional forward links added in a randomised way. The data in the list does have to be ordered, however, as during searching the elements in the array - bits of the list can be skipped to reduce the searching space quickly (like in a binary search). This leads to having an efficiency on all CRUD operations of $O(\log_2 n)$, which is logarithmic random time.

The ideal SkipList would have half of the nodes having one reference to a subsequent node; $\frac{1}{4}$ would have 2 references to subsequent nodes; $\frac{1}{8}$ will have 3 references to subsequent nodes; and so on. The distance between nodes on each 'level' would be equal.

SkipLists aren't guaranteed to give good performance. Their performance depends on how many nodes are skipped at each level. For searching, the best case will be $O(1)$ and the worst case will be $O(n)$.

Page 7

Async lecture - Hierarchical Data Structures

📅 2023-11-02



7.1 Why Another Data Structure?

The highest priority of any programmer when thinking about handling data is to ensure that the data is structured in a way to ensure that all the required operations to be performed on the data are efficient.

We have already learnt about a number of data structures: Static arrays which are fixed size; linked list which are dynamic sized; and SkipList which is also dynamic. The Static Array and SkipLists can be searched using a binary search permitting they are sorted however the Linked List can only be searched in a sequential fashion due to it's nature.

7.2 Hierarchical Data Structures

A hierarchical data structure (commonly a tree) consists of a collection of nodes where each node has a unique predecessor and many successors.

There is some key terminology to be aware of when discussing Hierarchical Data Structures, this is outlined in the table below and makes use of the following example of a tree.




Figure 7.1: Example of a Hierarchical Data Structure

Terminology	Definition	Example
Tree	A set of interconnected nodes with no closed circuits or loops	
Subtree	A tree rooted an internal node of the tree	Tree rooted at node B is a subtree of the given tree
Leaf (terminal node)	Any node of the tree which has no subtrees	The nodes containing D, H, F, I and J
Root node	The first node of the tree	The root node of the above tree is A
Branch	A link between two nodes within the tree	
Degree of a node	Number of subtrees of that node	The degree of the node containing A is 2; and B is 3
Level of a node	Number of branches on the path from the root to the node	The level of nodes B and C is 1; H, I and J is 3; A is 0
Height / Depth of a tree	Number of nodes on the longest path from the root node to a leaf node	The above tree has height 4; height of a tree containing one node is 1.

Table 7.1: Efficiency of CRUD operations for a static array

7.3 Binary Trees and Binary Search Trees

Binary Trees (BT) are a special case of tree where every node is of degree two or less. A Binary Search Tree (BST) is a special case of Binary tree where: the keys in the left subtree of the root precede the key in the root; the key in the root node precedes the keys in the right subtree; and the left and right subtrees of the root are also BSTs.

7.3.1 Searching

The algorithm to search a BST is based on the normal Binary Search algorithm (that we've used for a static array).

The first item to be evaluated is the root node, if this is the required item then we can stop our search. If the required item comes before the data item in the root node then we search the left subtree; if the required item comes after the data item stored in the root node then we search the right subtree. This process continues until either we find the required item or we reach a leaf node (hence the item is not in the BST).

7.3.2 Constructing a Binary Search Tree

When constructing a BST from scratch, you take the first data item you've been given and use that as the root node. From there add each data item as a child node, ensuring to maintain order.

7.3.3 Special Names for cases of Binary Trees

A binary tree is full if every node other than the leaf nodes has two children.

A binary tree is complete if all levels, except possibly the last level, are completely full and the last level has its node on the left side.

A balanced tree is one in which all path from the root node to the leaf nodes are of the same length.

7.3.4 Deletion of a node from a BST

In the case that the node we want to delete is a leaf node, we simply delete it.

In the case that the node has a single subtree then we simply replace the node to be deleted with the root node of the subtree.

In the case that the node to be deleted has two subtrees, we do something. Dalin hasn't actually told us what yet as that would be far too simple.

7.3.5 Traversal of Binary Trees

To traverse a data structure requires us to be able to access each node of the data structure once and only once in a predetermined sequence. Traversal of a tree is considerably more complex than that of a Static Array or even a Singly Linked List. There are two distinct approaches which can be taken to traverse a binary tree: DFT and BFT. Within each approach, there are sub-approaches - each of these will be explored below.

All examples shown below will make use of the following tree




Figure 7.2: Binary Search Tree used in traversal examples

7.3.5.1 Depth First Traversal

Depth First Travel (DFT) proceeds along a path from the root to the most distant descendent of the first child passed through before processing the second child. A stack is used to implement this.

The Depth First Traversal operates on the basis that we need to do three things: Visit the node (V), traverse the Left subtree (L), and traverse the Right subtree (R). There are three orders we do this in: VLR (PreOrder Traversal), LVR (InOrder Traversal) and LRV (PostOrder Traversal).

PreOrder Traversal In PreOrder Traversal, the binary tree is traversed as follows:

1. Visit the root node
2. PreOrder traversal of the left subtree
3. PreOrder traversal of the right subtree

Using the above example tree, a PreOrder Traversal would result in the following route: 9, 4, 2, 8, 13, 10, 25

InOrder Traversal In InOrder Traversal, the binary tree is traversed as follows:

1. InOrder Traversal of the left subtree
2. Visit the root node
3. InOrder traversal of the right subtree

Using the above example tree, an InOrder Traversal would result in the following route: 2, 4, 8, 9, 10, 13, 25

PostOrder Traversal In PostOrder Traversal, the binary tree is traversed as follows:

1. PostOrder Traversal of the left subtree
2. PostOrder Traversal of the right subtree
3. Visit the root node

Using the above example tree, a PostOrder Traversal would result in the following route: 2, 8, 4, 10, 25, 13, 9

7.3.5.2 Breadth First Traversal

Breadth First Traversal (BFT) proceeds horizontally from the root to all of its children then to its children's children and so on until all nodes have been processed. To implement this traversal, a queue is used.

Using the above example tree, a Breadth First Traversal would result in the following route: 9, 4, 13, 2, 8, 10, 25.

Page 8

Async lecture - Self Balancing Trees

 2023-11-09





8.1 Complexity of Binary Search Trees

When discussing the ideal case in a binary search tree - the tree will be balanced, or nearly balanced, as this minimises the height of the tree. This is important as the height of the tree relates to the number of comparisons we have to make whereby we make one comparison on each level. The worst case refers to the situation where there is only one node per level.

Operations	Ideal Case	Worst Case
Searching	$O(\log n)$	$O(n)$
Insertion	$O(\log n)$	$O(n)$
Deletion	$O(\log n)$	$O(n)$

Table 8.1: Complexity of a Binary Search Tree

8.2 Self Balancing Tree

A *self balancing tree* automatically keeps its height as small as possible at all times. It does this by performing “rotations” which reduce the height. These rotations change the structure of the tree without changing the order of elements therefore it remains a binary search tree.

8.2.1 Tree Rotation




Figure 8.1: Example of a rotation of a tree

Tree rotation is used to change the structure of the binary search tree without changing the order of the elements, thus after a rotation the tree is still a valid BST.

There are two popular approaches to implementing a self balancing tree - AVL Tree and Red-Black Tree (the later is not covered in the scope of this module).

8.3 AVL Trees

The *AVL Tree* is a type of Binary Search Tree which is named after it's inventors (Adelson, Velskii, and Landis). A BST is a valid AVL tree where:

1. the heights of the Left Subtree (LST) and Right Subtree (RST) of the root node differ by at most 1; and
2. LST and RST are AVL trees (this means we define it recursively)

8.3.1 Balance Factors

Each node of an AVL Tree contains an additional variable called the balance factor. The balance factor of a node can be calculated as follows

$$balanceFactor = height_{LST} - height_{RST}$$

We can use the balance factor to determine if a node is *balanced* or *unbalanced*. If the balance factor is -1 , 0 or $+1$ then the node is balanced. Whereas if the node has a balance factor of -2 or $+2$ then the node is unbalanced; this will also require rotation to be carried out so that the tree can become balanced.

8.3.2 Creating an AVL Tree

Much the same as a BST or a BT, an AVL tree is created as the data items are inserted as nodes. Initially the tree is empty, this means it satisfies the rules for an AVL tree. As data items are inserted, the balance factor of each node will change (nodes with balance factor ± 1 may be changed to ± 2). Where the balance factor is ± 2 then the tree is unbalanced and as such a rotation must take place. Rotation will take place around the unbalanced node which is closest to the most recently inserted node.

8.3.2.1 Rotations of an AVL Tree

There are only four possible rotations which can take place within an AVL tree, two single and two double.

Single Rotations:

- Left-Left Rotation (LL)
- Right-Right Rotation (RR)

Double Rotations:

- Left-Right Rotation (LR)
- Right-Left Rotation (RL)

These rotations take place around the unbalanced node with a balance factor of ± 2 .

The types of rotations are demonstrated below through use of an example. In the diagrams, the left hand tree shows the pre-rotation and the right hand tree shows the post-rotation layout.

8.3.2.2 Left-Left Rotation

The unbalanced node and the left subtree of the unbalanced node are both left heavy. The balance of the unbalanced node is $+2$ and the balance of the left child of the unbalanced node is $+1$. The tree is rebalanced by completing a single rotation to the right around the unbalanced node.




Figure 8.2: AVL tree rotation (Left-Left)

In the above example, the node containing 3 has just been inserted into the AVL tree - causing the node containing 8 to be unbalanced, therefore we rotate to the right about 8.

The pseudocode for the left-left rotation is as follows:

```
temp = root.left
root.left = root.left.right
temp.right = root
root = temp
```

8.3.2.3 Right-Right Rotation

The unbalanced node and the right subtree of the unbalanced node are both right-heavy. The balance of the unbalanced node is -2 and the balance of the right child of the unbalanced node is -1 . The tree is rebalanced by completing a single rotation to the left around the unbalanced node.

The Right-Right rotation is the mirror operation to the Left-Left rotation.




Figure 8.3: AVL tree rotation (Right-Right)

In the above example, the node containing 7 has just been inserted into the AVL tree causing the node containing 3 to become unbalanced. Therefore we rotate to the left about 3.

The pseudocode for the right-right rotation is as follows:

```
temp = root.right
root.right = temp.left
temp.left = root
root = temp
```

8.3.3 Right-Left Rotation

The unbalanced node is right heavy and the right subtree of the unbalanced node is left heavy. The balance of the unbalanced node is -2 and the balance of the right child of the unbalanced nodes is $+1$. This tree cannot be rebalanced by a single rotation - two rotations are required: a right rotation around the child followed by a left rotation around the unbalanced node.




Figure 8.4: AVL tree rotation (Right-Left)

The first rotation rotates node C right around node B, this results in the middle diagram. The second rotation rotates node A around node C, this results in the final state of the tree.

The pseudocode for the right-left rotation is shown below:

```
// setup
t1 = root.right
t2 = t1.left
// rotation 1: right
t1.left = t2.right
t2.right = t1
// rotation 2: left
root.right = t2.left
t2.left = root
root = t2
```

8.3.4 Left-Right Rotation

The unbalanced node is left-heavy and the left subtree of the unbalanced node is right-heavy. The balance factor of the unbalanced node is +2 and the balance factor of the left child of the unbalanced node is -1. This tree cannot be rebalanced by a single rotation - two rotations are required: a left around the child, followed by a right rotation around the unbalanced node.




Figure 8.5: AVL tree rotation (Left-Right)

The first rotation rotates node C left around node B, resulting in the middle diagram. The second rotation rotates node A right around node C, resulting in the right diagram.

Page 9

Async Lecture - Self-Organising Trees

📅 2023-11-09



Another type of tree is a *splay tree*. These are self-organising binary trees which organise themselves such that recently accessed elements are always quick to access again.

9.1 Basic Operations

All basic operations (insertion, search, removal) are performed in $O(\log_2 n)$ time and they all include an extra step called *splaying*.

Searching:

1. Locate node
2. Splay the node to the root

Insertion:

1. Insert node
2. Splay inserted node to the root

Deletion

1. Locate the node to be deleted
2. Replace node to be deleted with its in-order predecessor or its in-order successor (same process as deletion in a BST)
3. Splay the parent of the removed node to the root

9.2 Splaying

Splaying is the movement of a node to the root, which is achieved through a series of tree rotations which are similar to those used in AVL trees. Splaying works from the bottom to the top (root) of the tree, which is often called bottom-up splaying. Splaying involves a series of double rotations, until the accessed node reaches either the root or the child of the root when a single rotation is performed.

No single access operation on the Splay tree is guaranteed to be efficient, this is by design. The worst case will be $O(n)$, and the average across a series of operations tends to be $O(\log_2 n)$.

9.3 Rotations

9.3.1 Single Rotations

Single rotations are used when the accessed node *S* is the child of the root node *P*. The effect of this rotation is to move the node up one level in the tree.

9.3.1.1 Zig Rotation

There are two variations of the *Zig* rotation.




Figure 9.1: Splay tree rotations (Zig)

Rotation (a) is the equivalent of the LL rotation used with AVL trees; rotation (b) is the equivalent of the RR rotation used with AVL trees.

9.3.2 Double Rotations

Double rotations involve the accessed node *S*, the parent *P* and the grandparent *G*. The effect of these rotations is to move the node up two levels in the tree.

9.3.2.1 ZigZag Rotation

There are two variations of the *zigzag* rotation.




Figure 9.2: Splay tree rotations (ZigZag)

Rotation (c) is the equivalent of the RL rotation used with AVL trees; rotation (d) is the equivalent of LR rotation used with AVL trees. The ZigZag rotations tend to make the trees more balanced, and they will reduce the height of the tree by 1.

9.3.2.2 ZigZig Rotation

There are two variants of the *ZigZig* rotation.

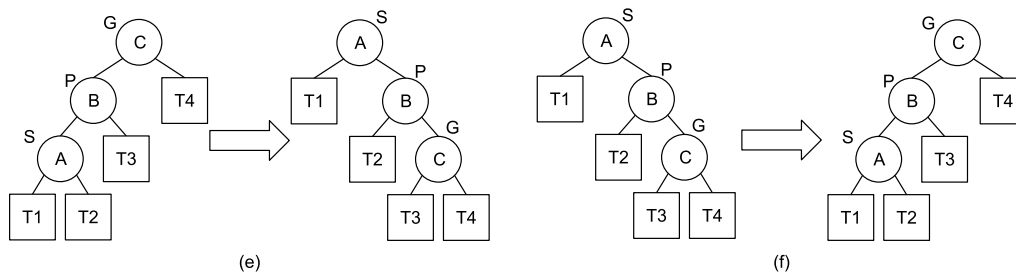


Figure 9.3: Splay tree rotations (ZigZig)

Pseudocode to the effect of rotation (e) is shown below:

```
G.left = P.right  
P.right = G
```

```
P.left = S.right  
S.right = P
```

Pseudocode to the effect of rotation (f) is shown below:

```
G.right = P.left  
P.left = G
```

```
P.right = S.left  
S.left = P
```